

Bayesian quantile correction and synthesis for climate products

Antonio Aguirre Raquel Prado Bruno Sansó
Department of Statistics, University of California, Santa Cruz • jaguir@ucsc.edu



QUESTION

Correction before synthesis

How can historical archives estimate source corrections before forecasts with different horizons are blended into one predictive distribution?

Case study: daily San Lorenzo River flow at the USGS Big Trees gauge, modeled on the $\log(1 + x)$ scale with x in m^3/s . Before each origin, USGS observations and retrospective ECMWF/GloFAS and NOAA/NWS products identify time-varying source corrections. After the origin, issued ensembles and exogenous information produce source-adjusted predictive flow quantiles.



USGS observations
15-min public gauge data; daily log-flow target for state updating and held-out verification.



ECMWF GloFAS
retrospective alignment; daily 51-member forecast guidance to 28 days.




NOAA NWS
retrospective alignment; hourly forecasts evaluated on common daily leads 1–8.

Exogenous information
Local basin precipitation and soil moisture; the first generalized dynamic principal component (GDPC) summarizes key Pacific, Atlantic, and global climate indices.

VALIDATION

Rolling-origin holdout design



Strict temporal holdout. At each origin, fitting uses only USGS flow and retrospective products available through that date. Issued forecasts and forecast-window exogenous inputs enter after the origin; future USGS observations are reserved for scoring. Main CRPS uses GloFAS-supported days 1–28, while NWS is compared over common leads 1–8.

INFERENCE

Scalable Bayesian quantile fitting

Seven independent Dynamic Quantile Linear Model (DQLM) lanes target seven levels: 0.05, 0.20, 0.35, 0.50, 0.65, 0.80, and 0.95. The fitted quantile forecasts are checked for crossing in the application and combined into the posterior predictive distribution scored by CRPS.

7

quantile lanes

≤12,995

USGS archive days


~2–4 h

fit + synthesis

- Parallel quantile lanes**
Each quantile level is fit as its own state-space model; the lanes can run concurrently before forecast quantiles are synthesized.
- Posterior approximation**
Mean-field variational Bayes (MFVB) retains closed-form Gaussian filtering and smoothing within latent-state blocks while approximating posterior dependence across blocks.
- Non-conjugate updates**
Variational Bayes with Laplace–Delta (VB-LD) uses local modes, curvature, and Delta-method moments for non-conjugate extended asymmetric Laplace (exAL) scale and skewness terms.
- State evolution**
Component discounting controls trend, multi-period seasonality, source corrections, and exogenous-covariate adjustments.

MODEL

One DQLM for source correction, then synthesis



DQLM correction stage. For a fixed quantile p_0 , one Dynamic Quantile Linear Model spans the pre-origin and forecast windows: USGS observations and retrospectives identify the latent flow quantile and source-specific discrepancies before T ; issued ensembles and staged covariates propagate forecast-window quantiles. Seven quantile lanes are then synthesized and scored only against held-out USGS observations.

Expanded quantile state-space model


For each target quantile p_0 and forecast origin T , the Dynamic Quantile Linear Model (DQLM) links observed flow, retrospective product histories, and issued forecast members through one latent river-flow quantile. Product-family discrepancies learn source corrections before the origin; the same source channels, with forecast-window covariates $\mathbf{x}_{T+k}$, then define corrected predictive quantiles after the origin.

Observed measurements:	y_t^o	$\sim \text{exAL}_{p_0}(\mathbf{F}_t^o \boldsymbol{\theta}_t + \zeta_t, \sigma^o, \gamma^o),$	$1 \leq t \leq T,$
Retrospective products:	z_t^j	$\sim \text{exAL}_{p_0}(\mathbf{F}_t^j(\boldsymbol{\theta}_t + \boldsymbol{\delta}_t^j) + \zeta_t, \sigma^j, \gamma^j),$	$j \in \mathcal{J}, 1 \leq t \leq T,$
Forecast products:	$y_{T+k}^j(k)$	$\sim \text{exAL}_{p_0}(\mathbf{F}_{T+k}^j(\boldsymbol{\theta}_{T+k} + \boldsymbol{\delta}_{T+k}^j) + \zeta_{T+k}, \sigma^j, \gamma^j),$	$j \in \mathcal{J}, i \in \mathcal{I}_j, k \in \mathcal{K}_j,$
Trend and seasonal states:	$\boldsymbol{\theta}_t \mid \boldsymbol{\theta}_{t-1}$	$\sim \mathcal{N}(\mathbf{G}_t \boldsymbol{\theta}_{t-1}, \mathbf{W}_t^{\boldsymbol{\theta}}),$	$1 \leq t \leq T + K_{\max},$
Source correction states:	$\boldsymbol{\delta}_t^j \mid \boldsymbol{\delta}_{t-1}^j$	$\sim \mathcal{N}(\mathbf{G}_t \boldsymbol{\delta}_{t-1}^j, \mathbf{W}_t^{\boldsymbol{\delta}^j}),$	$j \in \mathcal{J}, 1 \leq t \leq T + K_{\max},$
Exogenous adjustment:	$\zeta_t \mid \zeta_{t-1}, \psi_{t-1}$	$\sim \mathcal{N}(\lambda \zeta_{t-1} + \mathbf{x}_t^{\zeta} \psi_{t-1}, w_{\zeta}^2),$	$1 \leq t \leq T + K_{\max},$
Covariate-adjustment states:	$\psi_t \mid \psi_{t-1}$	$\sim \mathcal{N}(\psi_{t-1}, \mathbf{W}_t^{\psi}),$	$1 \leq t \leq T + K_{\max}.$

exAL_{p_0} is the extended asymmetric Laplace working likelihood for target quantile p_0 ; σ and γ are source-specific scale and skewness terms. T is the forecast origin, j a product family, i an ensemble member, k a lead, and $K_{\max}$ the longest horizon. The shared quantile state is $\boldsymbol{\theta}_t$; source discrepancies/corrections are $\boldsymbol{\delta}_t^j$; and (ζ_t, ψ_t) is the exogenous-adjustment block driven by covariates $\mathbf{x}_t$, such as precipitation, soil moisture, and the GDPC climate-index summary. Component discounting controls the evolution variances.

MAIN EVIDENCE


28-day CRPS comparison



Values left of the raw GloFAS line indicate lower mean CRPS than raw GloFAS at that origin; raw NWS is evaluated only in the separate 1–8 day panel.

HORIZON


Horizon-matched NWS check



Horizon-matched evidence. Over common days 1–8, exAL-M-T1 has the lowest mean CRPS at three origins and raw NWS at two. The axis is normalized by the origin-specific best 8-day CRPS, so 1.0 marks the winner for that origin; this is separate from the 28-day benchmark.

EXAMPLE


Representative forecast window



Illustrative, not aggregate validation. The selected exAL-M-T1 forecast at the 2022-12-25 origin shows how forecast inputs and posterior uncertainty evolve. At this origin AL-M-T1 has the lowest 28-day CRPS, so this panel is illustrative; aggregate evidence comes from all five origins.

DYNAMICS

80-month latent component



Retrospective diagnostic. From the selected exAL-M-T1 fit at the 2022-12-25 origin. Curves show the 5th, 50th, and 95th target-quantile components over 2000–2022; shading marks dry and wet periods for hydrologic context.

INTERPRETATION

What the evidence supports

- The source-aware exAL model (exAL-M-T1) has the lowest mean 28-day CRPS at four origins; AL-M-T1 is lowest at 2022-12-25.
- Over common days 1–8, exAL-M-T1 is lowest at three origins and raw NWS at two.
- Fitted latent states provide retrospective component diagnostics in addition to forecast-window predictive quantiles.
- The five-window design is a targeted stress test, not a continuous hindcast.

Take-home. Retrospective archives identify source-specific corrections before forecast-product blending. In this five-origin study, the source-aware exAL model leads the 28-day CRPS comparison at four origins and yields interpretable retrospective diagnostics.

