

Bayesian quantile-based correction and synthesis for climate products

Antonio Aguirre Raquel Prado Bruno Sansó
Department of Statistics, University of California, Santa Cruz • jaguir@ucsc.edu

CONTEXT

Archives and ensembles

Recent climate analysis and forecasting rely heavily on climate-center model archives and ensemble prediction systems. **Retrospective products** replay a center's physical–numerical modeling system over past conditions, so simulated flows can be compared with observations and source bias can be learned. **Forecast ensembles** run the operational system from current initial conditions to issue many plausible futures; their spread is useful, but biases, horizons, and update cycles differ by source.

QUESTION

Correct first, then synthesize

Can quantile-level bias be learned from diverse climate-product archives and propagated into newly issued ensemble forecasts on operational update timescales, while preserving an interpretable dynamic decomposition of source corrections, trend, seasonality, and exogenous effects?


Case study: daily San Lorenzo River flow at the USGS Big Trees gauge on the $\log(1 + x)$ scale. Before each forecast origin, observed flow is paired with retrospective ECMWF/GloFAS and NOAA/NWS products to estimate source corrections. After the origin, issued ensembles and staged exogenous covariates are used to carry those corrections into source-adjusted quantile forecasts and one posterior predictive distribution.



USGS observations
15-min public gauge record aggregated to daily log-flow; state-updating and held-out verification target.



ECMWF / GloFAS
European Centre for Medium-Range Weather Forecasts; retrospectives; daily 51-member ensemble forecasts, leads 1–28.




NOAA / NWS
U.S. National Weather Service under NOAA; retrospectives; hourly 7-member ensemble forecasts, daily leads 1–8.

Exogenous information
Local basin precipitation and soil moisture; the first generalized dynamic principal component (GDPC) summarizes key Pacific, Atlantic, and global climate indices.

VALIDATION

Rolling-origin holdout design



Strict temporal holdout. At each origin, fitting uses only USGS flow and retrospective products available through that date. Issued forecasts and forecast-window exogenous inputs enter after the origin; future USGS observations are reserved for scoring. Main CRPS uses GloFAS-supported days 1–28, while NWS is compared over common leads 1–8.

INFERENCE

Scalable Bayesian quantile fitting

Seven independent extended Dynamic Quantile Linear Model (exDQLM) lanes target seven levels: 0.05, 0.20, 0.35, 0.50, 0.65, 0.80, and 0.95. Posterior predictive samples are generated for each fitted quantile lane and synthesized into one posterior predictive distribution.

7
quantile lanes

≤12,995
USGS archive days


~2–4 h
fit + synthesis

Parallel quantile lanes
Each quantile level is fit as its own state-space model; the lanes can run concurrently before forecast quantiles are synthesized.

Scalable VB updates
Mean-field variational Bayes supports Kalman filtering and backward smoothing for state-distribution updates. Laplace–Delta approxima-

MODEL

Main Workflow



Source-aware exDQLM workflow. For each target quantile p_0 , pre-origin USGS observations and retrospective product histories identify a latent river-flow quantile and source-specific correction states. After the cutoff, issued ensemble forecasts and staged exogenous covariates propagate corrected forecast-window quantiles. The seven fitted quantile lanes are then synthesized into one posterior predictive distribution using isotonic adjustment and monotone rearranged draws, then scored only with held-out USGS observations.

extended Dynamic Quantile Linear Model (exDQLM)

For each target quantile p_0 , let T denote the forecast origin/cutoff. The selected exDQLM links observed flow, retrospective product histories, and issued forecast members through one latent river-flow quantile. Source discrepancies learn pre-origin corrections; forecast ensembles and staged covariates then propagate corrected quantiles after the origin.


Observed flow: $y_t^p \sim \text{exAL}_{p_0}(\mathbf{F}'_t \boldsymbol{\theta}_t + \zeta_t, \sigma^p, \gamma^p)$, $1 \leq t \leq T$,
Retrospectives: $z_t^j \sim \text{exAL}_{p_0}(\mathbf{F}'_t(\boldsymbol{\theta}_t + \boldsymbol{\delta}_t^j) + \zeta_t, \sigma^j, \gamma^j)$, $j \in \mathcal{J}$, $1 \leq t \leq T$,
Forecast ensembles: $y_{T+k}^{j,i} \sim \text{exAL}_{p_0}(\mathbf{F}'_{T+k}(\boldsymbol{\theta}_{T+k} + \boldsymbol{\delta}_{T+k}^j) + \zeta_{T+k}, \sigma^j, \gamma^j)$, $j \in \mathcal{J}$, $i \in \mathcal{I}_j$, $k \in \mathcal{K}_j$.

Trend/seasonal: $\boldsymbol{\theta}_t | \boldsymbol{\theta}_{t-1} \sim \mathcal{N}(\mathbf{G}_t \boldsymbol{\theta}_{t-1}, \mathbf{W}_t^{\boldsymbol{\theta}})$,
Discrepancies: $\boldsymbol{\delta}_t^j | \boldsymbol{\delta}_{t-1}^j \sim \mathcal{N}(\mathbf{G}_t \boldsymbol{\delta}_{t-1}^j, \mathbf{W}_t^{\boldsymbol{\delta}^j})$,
Transfer: $\zeta_t | \zeta_{t-1}, \boldsymbol{\psi}_{t-1} \sim \mathcal{N}(\lambda \zeta_{t-1} + \mathbf{x}'_t \boldsymbol{\psi}_{t-1}, w_t^\zeta)$,
Covariate effects: $\boldsymbol{\psi}_t | \boldsymbol{\psi}_{t-1} \sim \mathcal{N}(\boldsymbol{\psi}_{t-1}, \mathbf{W}_t^{\boldsymbol{\psi}})$.

Here $\boldsymbol{\theta}_t$ is the shared trend–seasonal river quantile, $\boldsymbol{\delta}_t^j$ is the source-specific correction, and $(\zeta_t, \boldsymbol{\psi}_t)$ is the exogenous-covariate block driven by $\mathbf{x}_t$ (local precipitation, local soil moisture, and GDPC climate summaries). **Priors:** $\boldsymbol{\theta}_0 \sim \mathcal{N}(\mathbf{m}_0, \mathbf{C}_0)$, $\sigma^p \sim IG(a_p, b_p)$, and $\gamma^p \sim t_{(L,U)}(0, \phi^p, \nu^p)$; component discounting controls $\mathbf{W}_t$, and time adaptation. DQLM uses conjugate AL; selected exDQLM uses exAL and VB-LD for non-conjugate σ, γ .


RESULTS

1–28-step-ahead forecast CRPS comparison



CRPS is the continuous ranked probability score; lower values are better. DQLM is the conjugate AL-likelihood baseline, while selected exDQLM denotes the exAL extension. Raw NWS is evaluated separately over common leads 1–8, not mixed into this 1–28 lead GloFAS-supported benchmark.


Wet-period quantile fit



Historical fit diagnostic. Selected exDQLM fitted quantile locations during the 2017–2019 wet period. The 5th, 50th, and 95th quantile fits are shown with observed USGS flow; wider upper-tail bands make winter high-flow behavior visible.

EXAMPLE


Three forecast-window examples



Origin-level illustrations. Three selected exDQLM forecast windows show how corrected quantile forecasts combine into a posterior predictive envelope after the cutoff. The panels are illustrative checks of behavior at individual origins; the aggregate comparison remains the five-origin CRPS benchmark.

DYNAMICS

Interpretable seasonal component



Retrospective diagnostic. One component of the interpretable state decomposition. From the selected exDQLM at the 2022-12-25 origin, lines show the 80-month seasonal state for the 5th, 50th, and 95th target-quantile lanes over 2008–2022; light ribbons show 95% credible bands. Shaded dry (2012–2016) and wet (2017–2019) regimes mark Santa Cruz-relevant water-supply contrasts, and the recovered long-cycle state tracks this multi-year behavior.

TAKEAWAYS

Main takeaways

- Source correction.** Quantile-level bias is learned from retrospective archives and propagated into newly issued ensemble forecasts.
- Operational synthesis.** Seven corrected quantile lanes are fit in parallel and synthesized into one posterior predictive distribution on update timescales.
- Interpretable dynamics.** The state decomposition separates source corrections from trend, seasonality, and exogenous effects, linking product bias to river-flow dynamics.

Limitations and next steps

- Quantile coherence.** Crossing was not material here; joint quantile estimation could improve coherence and possibly forecast skill.
- Nonlinear response.** Linear state dynamics aid interpretation, but may miss nonlinear hydrologic behavior.
- Deployment tuning.** Discount factors and transfer rate λ require tuning or model selection before operational use.

Code and reproducibility
github.com/AntonioAPDL/Project1

