ISBA 2026
World Meeting, Nagoya, Japan

Bayesian quantile correction and synthesis for climate products

Antonio Aguirre   Raquel Prado   Bruno Sansó
Department of Statistics, University of California, Santa Cruz • jaguir@ucsc.edu

QUESTION

Fast correction before synthesis

Can historical product archives learn source discrepancies, propagate those corrections into new quantile-level ensemble forecasts, and synthesize calibrated posterior predictions fast enough for real-time product updates?

Case study: daily San Lorenzo River flow at the USGS Big Trees gauge on the $\log(1 + x)$ scale. Before each origin, USGS flow and retrospective GloFAS/NWS products learn dynamic source discrepancies relative to the gauge. After the origin, issued ensembles and exogenous covariates carry those corrections into source-adjusted quantile forecasts and one posterior predictive flow distribution.

USGS observations
15-min public gauge data; daily log-flow target for state updating and held-out verification.

ECMWF GloFAS
retrospective alignment; daily 51-member forecast guidance to 28 days.

NOAA NWS
retrospective alignment; hourly forecasts evaluated on common daily leads 1–8.

Exogenous information
Local basin precipitation and soil moisture; the first generalized dynamic principal component (GDPC) summarizes key Pacific, Atlantic, and global climate indices.

VALIDATION

Rolling-origin holdout design

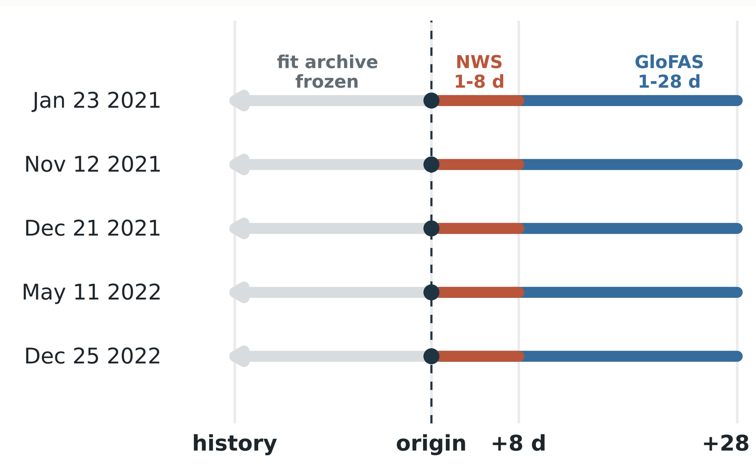

Strict temporal holdout. At each origin, fitting uses only USGS flow and retrospective products available through that date. Issued forecasts and forecast-window exogenous inputs enter after the origin; future USGS observations are reserved for scoring. Main CRPS uses GloFAS-supported days 1–28, while NWS is compared over common leads 1–8.

INFERENCE

Scalable Bayesian quantile fitting

Seven independent extended Dynamic Quantile Linear Model (exDQLM) lanes target seven levels: 0.05, 0.20, 0.35, 0.50, 0.65, 0.80, and 0.95. The fitted quantile forecasts are checked for crossing in the application and combined into the posterior predictive distribution scored by CRPS.

7	≤12,995	∼2–4 h
quantile lanes	USGS archive days	fit + synthesis

Parallel quantile lanes
Each quantile level is fit as its own state-space model; the lanes can run concurrently before forecast quantiles are synthesized.

Conjugate DQLM baseline
The asymmetric-Laplace DQLM is fully conjugate under its latent representation; this is the baseline labeled DQLM in the Results panel.

Non-conjugate exDQLM updates
The selected exDQLM replaces AL by exAL. Mean-field variational Bayes keeps Kalman filtering and backward smoothing for Gaussian state blocks, while VB-LD uses local modes, curvature, and Delta-method moments for non-conjugate exAL scale and skewness terms.

State evolution
Component discounting controls trend, multi-period seasonality, source corrections, and exogenous-covariate adjustments.

MODEL

Main Workflow

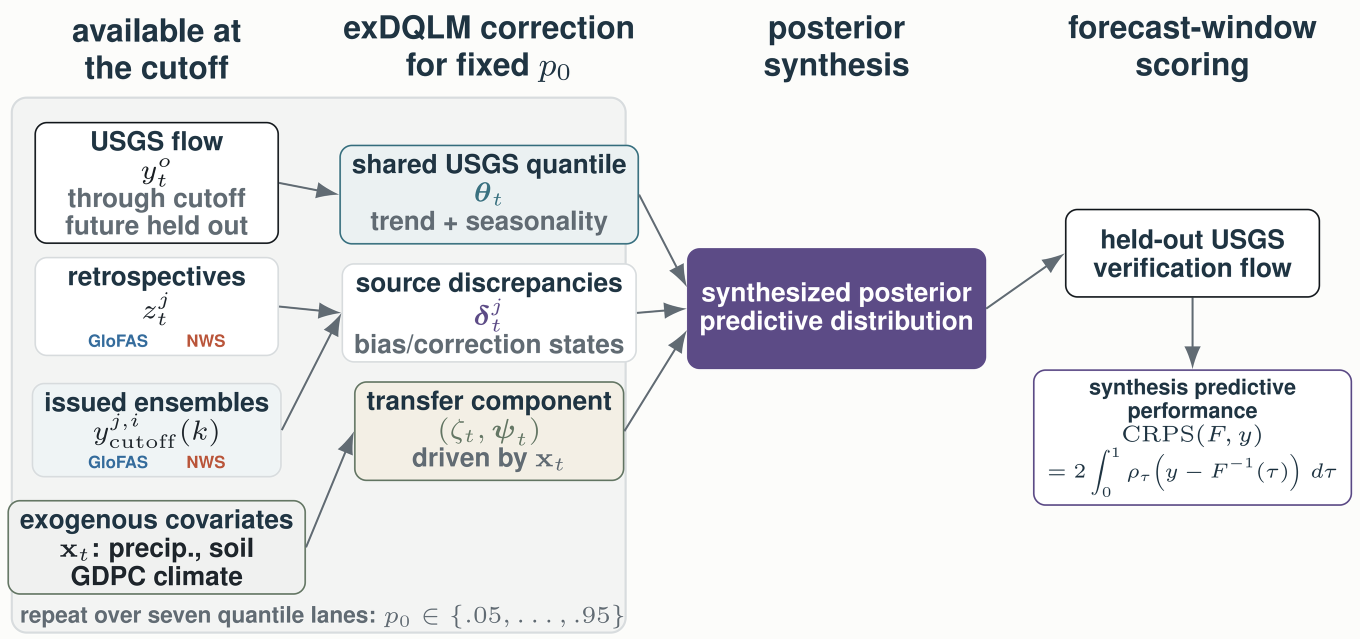

Source-aware exDQLM workflow. For each target quantile p_0, pre-origin USGS observations and retrospective product histories identify a latent river-flow quantile and source-specific correction states. After the cutoff, issued ensemble forecasts and staged exogenous covariates propagate corrected forecast-window quantiles. The seven fitted quantile lanes are then synthesized into one posterior predictive distribution and scored only with held-out USGS observations.

extended Dynamic Quantile Linear Model

For each target quantile p_0, let T denote the forecast origin/cutoff. The extended Dynamic Quantile Linear Model (exDQLM) links observed flow, retrospective product histories, and issued forecast members through one latent river-flow quantile. Product-family discrepancies learn source corrections before the origin; the same source channels, with forecast-window covariates $\mathbf{x}_{T+k}$, then define corrected predictive quantiles after the origin.

$$
\begin{array}{rlll}
\text{Observed measurements:} & y_t^o & \sim \mathrm{exAL}_{p_0}(\mathbf{F}_t'\boldsymbol{\theta}_t + \zeta_t, \sigma^o, \gamma^o)\,, & 1 \le t \le T, \\
\text{Retrospective products:} & z_t^j & \sim \mathrm{exAL}_{p_0}\left(\mathbf{F}_t'(\boldsymbol{\theta}_t + \boldsymbol{\delta}_t^j) + \zeta_t, \sigma^j, \gamma^j\right), & j \in \mathcal{J},\ 1 \le t \le T, \\
\text{Forecast products:} & y_T^{j,i}(k) & \sim \mathrm{exAL}_{p_0}\left(\mathbf{F}_{T+k}'(\boldsymbol{\theta}_{T+k} + \boldsymbol{\delta}_{T+k}^j) + \zeta_{T+k}, \sigma^j, \gamma^j\right), & j \in \mathcal{J},\ i \in \mathcal{I}_j,\ k \in \mathcal{K}_j, \\
\text{Trend and seasonal states:} & \boldsymbol{\theta}_t \mid \boldsymbol{\theta}_{t-1} & \sim \mathcal{N}\left(\mathbf{G}_t\boldsymbol{\theta}_{t-1}, \mathbf{W}_t^{\theta}\right), & 1 \le t \le T + K_{\max}, \\
\text{Source correction states:} & \boldsymbol{\delta}_t^j \mid \boldsymbol{\delta}_{t-1}^j & \sim \mathcal{N}\left(\mathbf{G}_t\boldsymbol{\delta}_{t-1}^j, \mathbf{W}_t^{\delta^j}\right), & j \in \mathcal{J},\ 1 \le t \le T + K_{\max}, \\
\text{Exogenous adjustment:} & \zeta_t \mid \zeta_{t-1}, \boldsymbol{\psi}_{t-1} & \sim \mathcal{N}\left(\lambda\zeta_{t-1} + \mathbf{x}_t'\boldsymbol{\psi}_{t-1}, w_t^{\zeta}\right), & 1 \le t \le T + K_{\max}, \\
\text{Covariate-adjustment states:} & \boldsymbol{\psi}_t \mid \boldsymbol{\psi}_{t-1} & \sim \mathcal{N}\left(\boldsymbol{\psi}_{t-1}, \mathbf{W}_t^{\psi}\right), & 1 \le t \le T + K_{\max}.
\end{array}
$$

exAL_{p_0} is the extended asymmetric Laplace working likelihood for target quantile p_0; σ and γ are source-specific scale and skewness terms. T is the forecast origin, j a product family, i an ensemble member, k a lead, and $K_{\max}$ the longest horizon. The shared quantile state is $\boldsymbol{\theta}_t$; source discrepancies/corrections are $\boldsymbol{\delta}_t^j$; and $(\zeta_t, \boldsymbol{\psi}_t)$ is the exogenous-adjustment block driven by covariates $\mathbf{x}_t$ such as precipitation, soil moisture, and the GDPC climate-index summary. Component discounting controls the evolution variances. The DQLM comparison in the Results panel is the conjugate AL-likelihood baseline; the selected exDQLM uses exAL and updates the non-conjugate σ, γ terms with VB-LD.

RESULTS

1–28-step-ahead forecast CRPS comparison

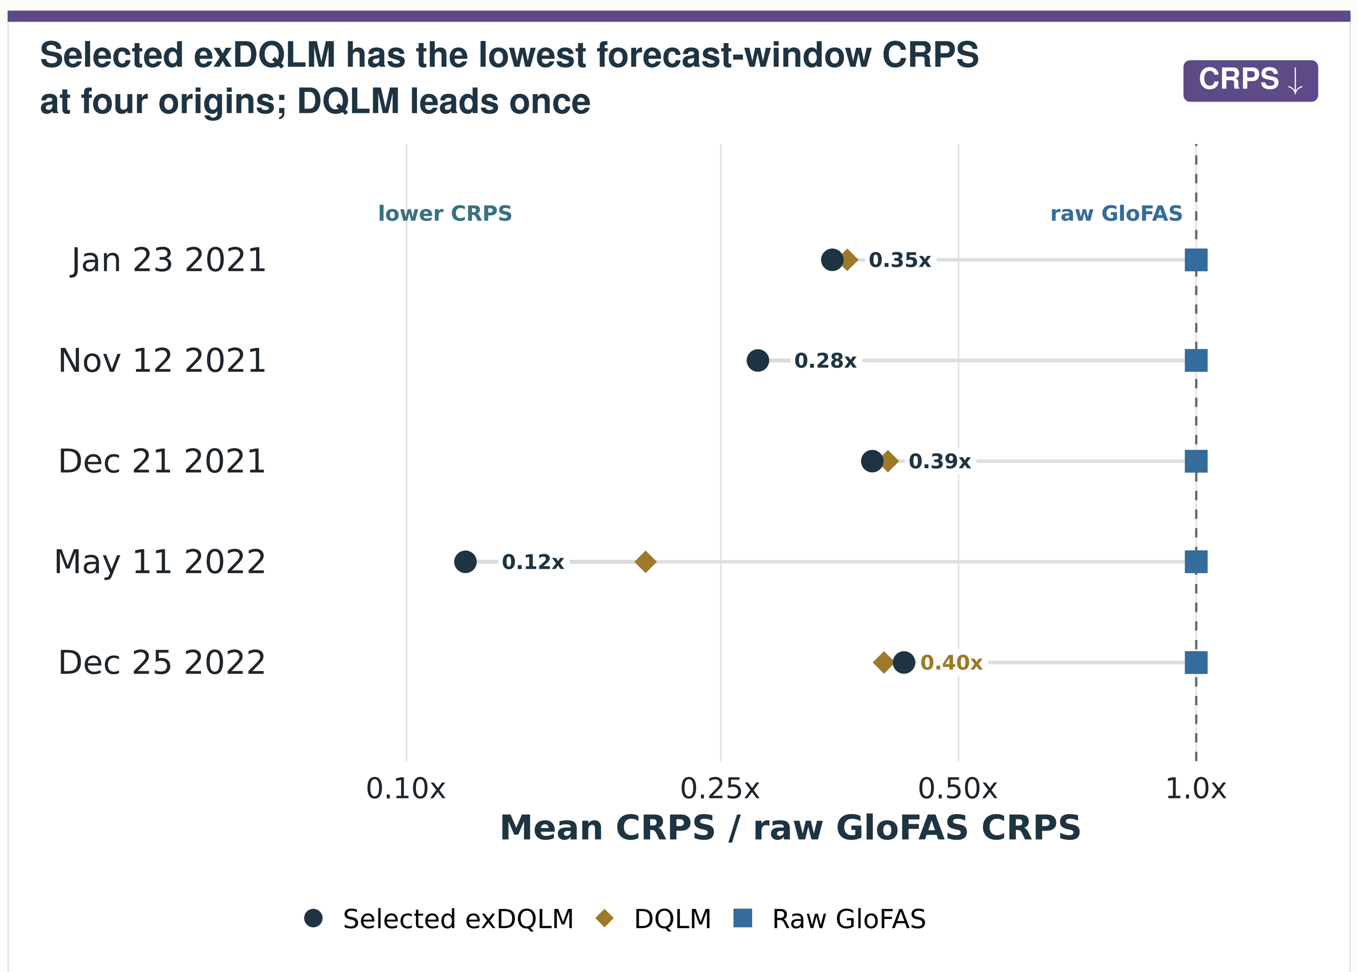

Values left of the raw GloFAS line indicate lower mean CRPS than raw GloFAS at that origin. DQLM denotes the conjugate AL-likelihood baseline; selected exDQLM denotes the exAL extension. Raw NWS is evaluated over common leads 1–8 in the horizon-matched tables, not mixed into this 1–28 lead GloFAS-supported benchmark.

FIT

Wet-period quantile fit

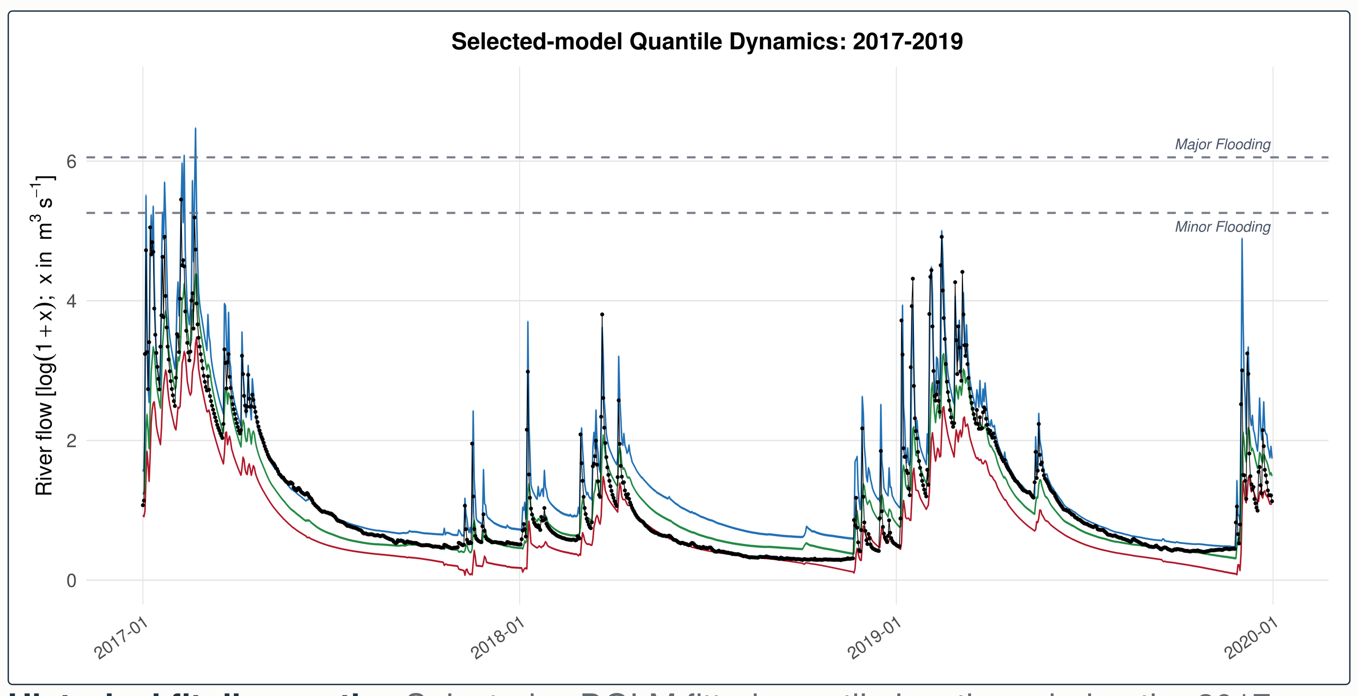

Historical fit diagnostic. Selected exDQLM fitted quantile locations during the 2017–2019 wet period. The 5th, 50th, and 95th quantile fits are shown with observed USGS flow; wider upper-tail bands make winter high-flow behavior visible.

EXAMPLE

Three forecast-window examples

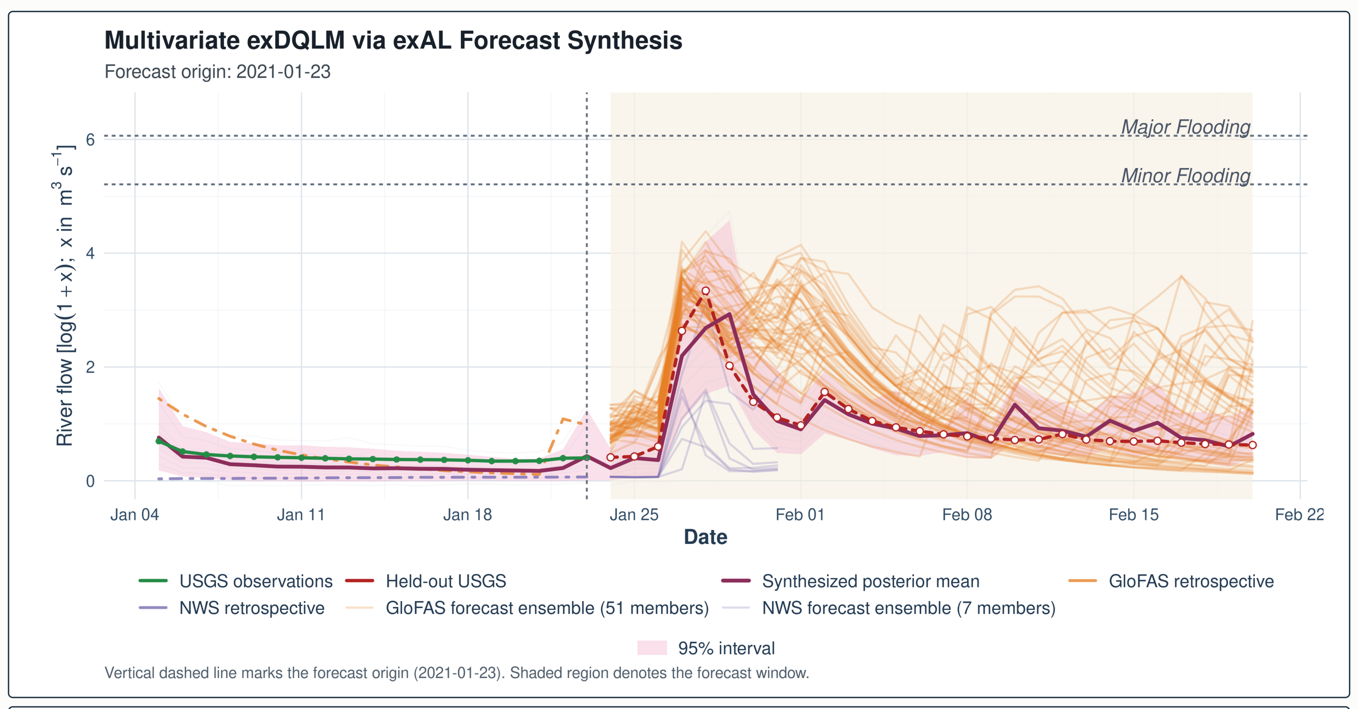

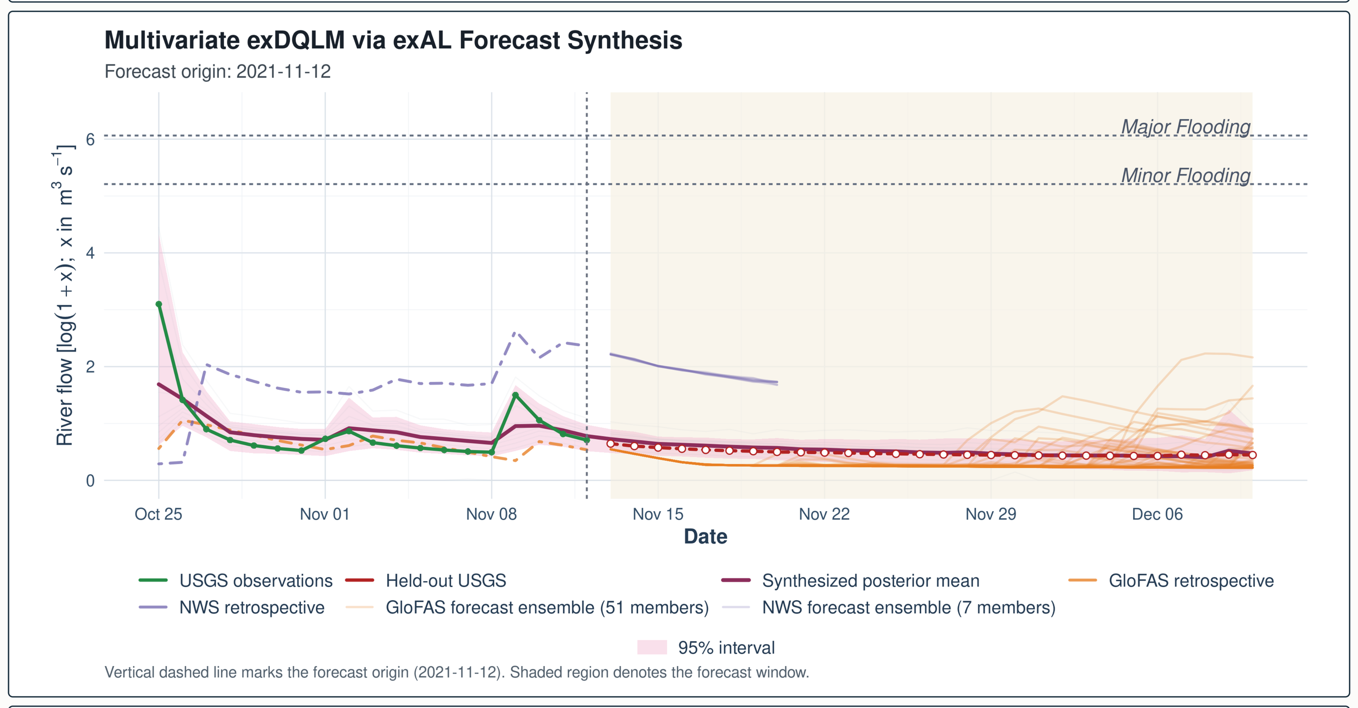

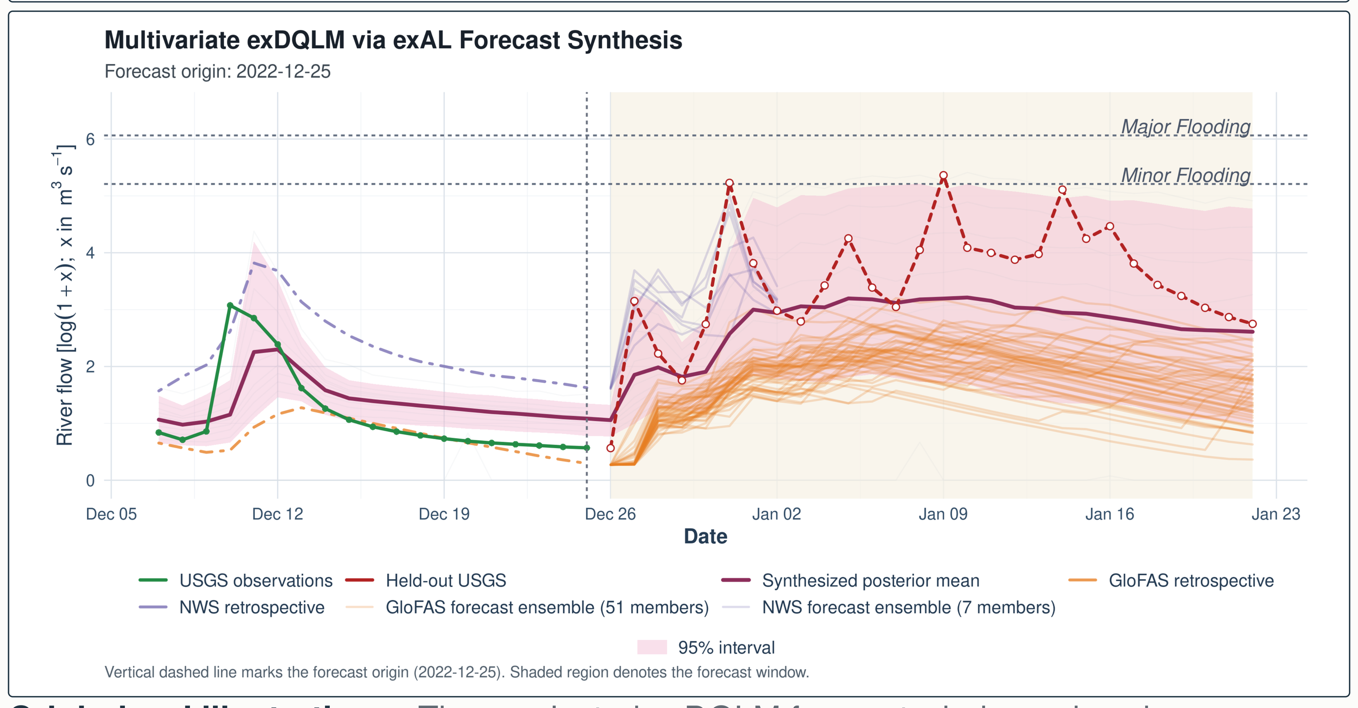

Origin-level illustrations. Three selected exDQLM forecast windows show how corrected quantile forecasts combine into a posterior predictive envelope after the cutoff. The panels are illustrative checks of behavior at individual origins; the aggregate comparison remains the five-origin CRPS benchmark.

DYNAMICS

80-month seasonal component

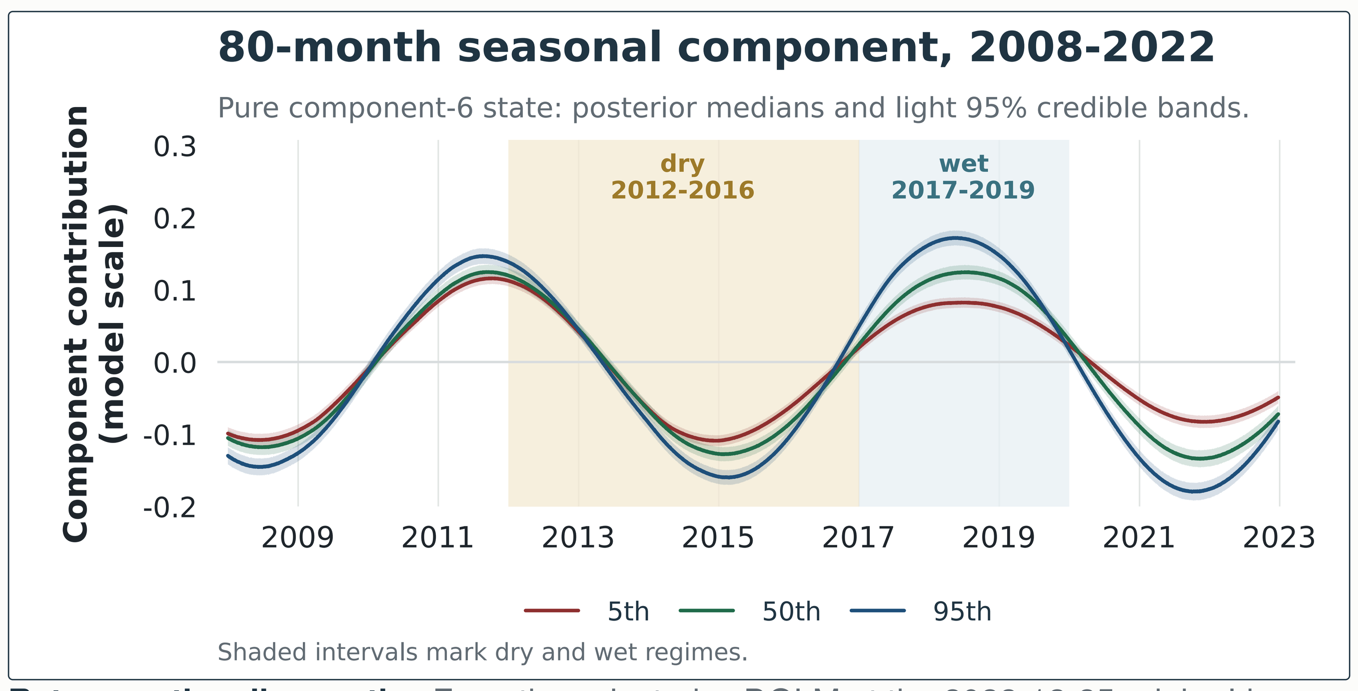

Retrospective diagnostic. From the selected exDQLM at the 2022-12-25 origin. Lines show the isolated 80-month state component, without the trend component added, for the 5th, 50th, and 95th target-quantile lanes over 2008–2022; light ribbons show 95% credible bands and shaded intervals mark dry and wet regimes.

TAKEAWAYS

Main takeaways

- Historical archives learn source-specific corrections before new forecast products are blended.
- The selected exDQLM extends the conjugate DQLM with exAL scale and skewness flexibility.
- Seven corrected quantile lanes synthesize one posterior predictive distribution.
- Selected exDQLM has the lowest 1–28-step-ahead CRPS at four of five origins; DQLM leads at the final origin.

Code and reproducibility
github.com/AntonioAPDL/Project1