

Bayesian quantile correction and synthesis for climate products

Antonio Aguirre Raquel Prado Bruno Sansó
Department of Statistics, University of California, Santa Cruz • jaguir@ucsc.edu



QUESTION

Fast synthesis of corrected products

How can large retrospective product archives be used efficiently at an operational forecast origin to learn bias for each source, calibrate forecast ensembles across lead times, and synthesize predictive quantiles for the target process?

Case study: daily San Lorenzo River flow at the USGS Big Trees gauge. The framework learns dynamic corrections from USGS observations and retrospective ECMWF/GloFAS and NOAA/NWS products, then carries those corrections into the forecast window to synthesize calibrated predictive flow quantiles. This is needed because product families differ in lead time, update cycle, ensemble structure, and uncertainty representation.



USGS observations
15-min public gauge data; daily flow used for state updating and held-out verification.



ECMWF / GloFAS
daily issues; retrospective flow for bias correction; 51-member forecasts to 28 days.




NOAA / NWS
hourly issues; retrospective flow for bias correction; forecasts on common daily leads 1–8.

Exogenous information
Local basin precipitation and soil moisture; the first generalized dynamic principal component (GDPC) summarizes key Pacific, Atlantic, and global climate indices.

MODEL

One state-space model for correction and synthesis

Source-aware quantile correction and synthesis
Retrospective products identify source corrections before the origin; forecast products enter through those corrected channels.



Correction before synthesis. Before the origin, observations and retrospective products identify the shared quantile state and source corrections. After the origin, issued forecast products and exogenous covariates update the same state to form calibrated predictive quantiles.

Expanded quantile state-space model
For each target quantile p_0 and forecast origin T , observed flow, retrospective product histories, and issued forecast members are linked through one latent river-flow quantile. Retrospective products identify source corrections before the origin; forecast products enter through those corrected source channels.

Observed measurements: $y_t^o \sim \text{exAL}_{p_0}(\mathbf{F}_t' \boldsymbol{\theta}_t + \zeta_t, \sigma^o, \gamma^o),$

Retrospective products: $z_t^j \sim \text{exAL}_{p_0}(\mathbf{F}_t'(\boldsymbol{\theta}_t + \boldsymbol{\delta}_t^j) + \zeta_t, \sigma^j, \gamma^j),$

Forecast products: $y_{T+k}^j(k) \sim \text{exAL}_{p_0}(\mathbf{F}_{T+k}'(\boldsymbol{\theta}_{T+k} + \boldsymbol{\delta}_{T+k}^j) + \zeta_{T+k}, \sigma^j, \gamma^j),$

Trend and seasonal states: $\boldsymbol{\theta}_t | \boldsymbol{\theta}_{t-1} \sim \mathcal{N}(\mathbf{G} \boldsymbol{\theta}_{t-1}, \mathbf{W}_t^{\boldsymbol{\theta}}),$

Source correction states: $\boldsymbol{\delta}_t^j | \boldsymbol{\delta}_{t-1}^j \sim \mathcal{N}(\mathbf{G}_j \boldsymbol{\delta}_{t-1}^j, \mathbf{W}_t^{\boldsymbol{\delta}_j}),$

Exogenous adjustment: $\zeta_t | \zeta_{t-1}, \boldsymbol{\psi}_{t-1} \sim \mathcal{N}(\lambda \zeta_{t-1} + \mathbf{x}_t' \boldsymbol{\psi}_{t-1}, w_t^{\zeta}),$

Covariate-effect states: $\boldsymbol{\psi}_t | \boldsymbol{\psi}_{t-1} \sim \mathcal{N}(\boldsymbol{\psi}_{t-1}, \mathbf{W}_t^{\boldsymbol{\psi}}),$

$1 \leq t \leq T,$

$j \in \mathcal{J}, 1 \leq t \leq T,$

$j \in \mathcal{J}, i \in \mathcal{I}_j, k \in \mathcal{K}_j,$

$1 \leq t \leq T + K_{\max},$

$j \in \mathcal{J}, 1 \leq t \leq T + K_{\max},$

$1 \leq t \leq T + K_{\max},$

$1 \leq t \leq T + K_{\max}.$


exAL_{p_0} is the extended asymmetric Laplace likelihood; σ and γ are source-specific scale and skewness terms. T is the forecast origin, j a product family, i an ensemble member, k a lead, $K_{\max}$ the longest horizon, and $\mathbf{x}_t$ the exogenous covariates. $\boldsymbol{\theta}_t$ is the shared quantile state, $\boldsymbol{\delta}_t^j$ is the source correction/discrepancy state, and $(\zeta_t, \boldsymbol{\psi}_t)$ is the covariate-effect block. Priors use bounded truncated Student- t terms for γ^j , inverse-gamma terms for σ^j , Gaussian initial states, and component discounting for evolution variances.

MAIN EVIDENCE

Held-out forecast-window distributional accuracy

28-day CRPS across five held-out rolling origins

Mean CRPS relative to raw GloFAS at the same cutoff; lower and farther left



● Selected model ◆ AL synthesis ■ GloFAS ● Other Bayesian variants

Grey points are the remaining Bayesian benchmark variants.

Use all five origins for final claims.


How to read this panel. Lower 28-day CRPS indicates better distributional accuracy over the held-out forecast window at each origin.

Values left of the raw GloFAS line indicate lower mean CRPS at the same origin.

HORIZON

Horizon-matched NWS check

8-day NWS-compatible comparison
Days 1-8 only; lower is better.



● Selected model ◆ AL synthesis ▲ NWS ■ GloFAS


This is a horizon-matched comparison, not the 28-day medium-range re

Horizon-matched evidence. NWS is available over common valid leads 1–8 in this evaluation. It is scored separately so the 28-day medium-range benchmark is not mixed with the NWS 1–8 day horizon.

EXAMPLE

Representative forecast window

Multivariate exDQLM via exAL Forecast Synthesis
Forecast origin: 2022-12-25




Illustrative, not aggregate validation. The panel shows how forecast inputs and posterior uncertainty evolve at the 2022-12-25 origin. It is illustrative; aggregate validation comes from the five-origin CRPS panel.

DYNAMICS

80-month component dynamics

80-month component, 2000-2022
Selected 2022-12-25 fit; posterior medians by target quantile.



More than forecast correction. The same state-space fit yields retrospective component diagnostics. The long-period harmonic tracks slow wet–dry structure while preserving quantile-specific variation.

INTERPRETATION

What the evidence supports

- Historically aligned retrospective products identify source corrections before the forecast window.
- Issued ensembles and exogenous covariates enter the state-space model directly, rather than being averaged after prediction.
- Observation-only quantile references have about 5–41x higher 28-day CRPS than the selected source-aware model across these origins.
- The five-window design is a targeted stress test, not a continuous hindcast; broader origin coverage is the next validation step.

Take-home. Learning source-specific corrections before synthesis can improve medium-range predictive distributions while preserving interpretable latent dynamics.

UC Santa Cruz Department of Statistics
Contact: jaguir@ucsc.edu • Data and product logos identify sources; no endorsement is implied. • Poster, manuscript, supplementary diagnostics, and reproducibility materials are available through the project repository.

Project repo Code and supplementary material
Project repository