


QUESTION


Correct first, then synthesize

Can quantile-level bias be learned from diverse climate-product archives and propagated into newly issued ensemble forecasts on operational update timescales, while preserving an interpretable dynamic decomposition of source corrections, trend, seasonality, and exogenous effects?

Case study: daily San Lorenzo River flow at the USGS Big Trees gauge on the $\log(1 + x)$ scale. Before each forecast origin, observed flow is paired with retrospective ECMWF/GloFAS and NOAA/NWS products to estimate source corrections. After the origin, issued ensembles and staged exogenous covariates are used to carry those corrections into source-adjusted quantile forecasts and one posterior predictive distribution.


USGS observations
15-min public gauge record aggregated to daily log-flow; state-updating and held-out verification target.


ECMWF / GloFAS
European Centre for Medium-Range Weather Forecasts; retrospectives; daily 51-member ensemble forecasts, leads 1–28.



NOAA / NWS
U.S. National Weather Service under NOAA; retrospectives; hourly 7-member ensemble forecasts, daily leads 1–8.

Exogenous information

Local basin precipitation and soil moisture; the first generalized dynamic principal component (GDPC) summarizes key Pacific, Atlantic, and global climate indices.

VALIDATION

Rolling-origin holdout design



Strict temporal holdout. At each origin, fitting uses only USGS flow and retrospective products available through that date. Issued forecasts and forecast-window exogenous inputs enter after the origin; future USGS observations are reserved for scoring. Main CRPS uses GloFAS-supported days 1–28, while NWS is compared over common leads 1–8.

INFERENCE

Scalable Bayesian quantile fitting

Seven independent extended Dynamic Quantile Linear Model (exDQLM) lanes target seven levels: 0.05, 0.20, 0.35, 0.50, 0.65, 0.80, and 0.95. Posterior predictive samples are generated for each fitted quantile lane and synthesized into one posterior predictive distribution.

7
quantile lanes

≤12,995
USGS archive days


~2–4 h
fit + synthesis

Parallel quantile lanes
Each quantile level is fit as its own state-space model; the lanes can run concurrently before forecast quantiles are synthesized.

Scalable VB updates
Mean-field variational Bayes supports Kalman filtering and backward smoothing for state-distribution updates. Laplace–Delta approximations handle the non-conjugate exAL scale and skewness parameters.

MODEL

Main Workflow



Source-aware exDQLM workflow. For each target quantile p_0 , pre-origin USGS observations and retrospective product histories identify a latent river-flow quantile and source-specific correction states. After the cutoff, issued ensemble forecasts and staged exogenous covariates propagate corrected forecast-window quantiles. The seven fitted quantile lanes are then synthesized into one posterior predictive distribution and scored only with held-out USGS observations.

Dynamic Quantile Linear Model (exDQLM)


For each target quantile p_0 , let T denote the forecast origin/cutoff. The selected exDQLM links observed flow, retrospective product histories, and issued forecast members through one latent river-flow quantile. Source discrepancies learn pre-origin corrections; forecast ensembles and staged covariates then propagate corrected quantiles after the origin.

$$y_t^o \sim \text{exAL}_{p_0}(\mathbf{F}_t' \boldsymbol{\theta}_t + \zeta_t, \sigma^o, \gamma^o), \quad 1 \leq t \leq T,$$
$$z_t^j \sim \text{exAL}_{p_0}(\mathbf{F}_t' \boldsymbol{\theta}_t + \boldsymbol{\delta}_t^j + \zeta_t, \sigma^j, \gamma^j), \quad j \in \mathcal{J}, 1 \leq t \leq T,$$
$$y_T^{j,i}(k) \sim \text{exAL}_{p_0}(\mathbf{F}_{T+k}'(\boldsymbol{\theta}_{T+k} + \boldsymbol{\delta}_{T+k}^j + \zeta_{T+k}, \sigma^j, \gamma^j), \quad j \in \mathcal{J}, i \in \mathcal{I}_j, k \in \mathcal{K}_j.$$
$$\boldsymbol{\theta}_t \mid \boldsymbol{\theta}_{t-1} \sim \mathcal{N}(\mathbf{G}_t \boldsymbol{\theta}_{t-1}, \mathbf{W}_t^\theta), \quad \boldsymbol{\delta}_t^j \mid \boldsymbol{\delta}_{t-1}^j \sim \mathcal{N}(\mathbf{G}_t \boldsymbol{\delta}_{t-1}^j, \mathbf{W}_t^{\delta^j}).$$
$$\zeta_t \mid \zeta_{t-1}, \boldsymbol{\psi}_{t-1} \sim \mathcal{N}(\lambda \zeta_{t-1} + \mathbf{x}_t' \boldsymbol{\psi}_{t-1}, w_\zeta^2), \quad \boldsymbol{\psi}_t \mid \boldsymbol{\psi}_{t-1} \sim \mathcal{N}(\boldsymbol{\psi}_{t-1}, \mathbf{W}_t^\psi).$$

Here $\boldsymbol{\theta}_t$ is the shared trend–seasonal river quantile, $\boldsymbol{\delta}_t^j$ is the source-specific correction, and $(\zeta_t, \boldsymbol{\psi}_t)$ is the exogenous-covariate block driven by $\mathbf{x}_t$ (local precipitation, local soil moisture, and GDPC climate summaries). The DQLM comparator uses the conjugate AL likelihood; the selected exDQLM uses exAL and updates non-conjugate σ, γ terms with VB-LD.

RESULTS


1–28-step-ahead forecast CRPS comparison



CRPS is the continuous ranked probability score; lower values are better. DQLM is the conjugate AL-likelihood baseline, while selected exDQLM denotes the exAL extension. Raw NWS is evaluated separately over common leads 1–8, not mixed into this 1–28 lead GloFAS-supported benchmark.

FIT


Wet-period quantile fit



Historical fit diagnostic. Selected exDQLM fitted quantile locations during the 2017–2019 wet period. The 5th, 50th, and 95th quantile fits are shown with observed USGS flow; wider upper-tail bands make winter high-flow behavior visible.

EXAMPLE


Three forecast-window examples



Origin-level illustrations. Three selected exDQLM forecast windows show how corrected quantile forecasts combine into a posterior predictive envelope after the cutoff. The panels are illustrative checks of behavior at individual origins; the aggregate comparison remains the five-origin CRPS benchmark.

DYNAMICS

Interpretable seasonal component



Retrospective diagnostic. One component of the interpretable state decomposition. From the selected exDQLM at the 2022-12-25 origin, lines show the isolated 80-month seasonal state, without the trend component added, for the 5th, 50th, and 95th target-quantile lanes over 2008–2022; light ribbons show 95% credible bands and shaded intervals mark dry and wet regimes.

TAKEAWAYS

Main takeaways

- Selected exDQLM has the lowest 1–28-step-ahead CRPS at four of five origins; DQLM leads at 2022-12-25.
- The model estimates quantile-level source bias from archived retrospectives, then propagates corrections into newly issued ensembles.
- Seven corrected quantile lanes synthesize one posterior predictive distribution on operational update timescales.
- The state decomposition separates source corrections, trend, seasonality, and exogenous effects; the 80-month panel shows one diagnostic component.

