

Bayesian quantile correction and synthesis for climate products

Antonio Aguirre Raquel Prado Bruno Sansó
Department of Statistics, University of California, Santa Cruz • jaguir@ucsc.edu



QUESTION

Correction before synthesis

How can historical archives estimate source corrections before forecasts with different horizons are blended into one predictive distribution?


Case study: daily San Lorenzo River flow at the USGS Big Trees gauge, modeled on the $\log(1 + x)$ scale with x in m^3/s . Before each forecast origin, USGS observations and retrospective ECMWF/-GloFAS and NOAA/NWS products identify time-varying source corrections. After that origin, issued ensembles and exogenous information produce source-adjusted predictive flow quantiles.



USGS observations
15-min public gauge data; daily log-flow target for state updating and held-out verification.



ECMWF GloFAS retrospective alignment; daily 51-member forecast guidance to 28 days.




NOAA NWS retrospective alignment; hourly forecasts evaluated on common daily leads 1–8.

Exogenous information

Local basin precipitation and soil moisture; the first generalized dynamic principal component (GDPC) summarizes key Pacific, Atlantic, and global climate indices.

VALIDATION

Rolling-origin holdout design



Strict temporal holdout. At each origin, fitting uses only USGS flow and retrospective products available through that date. Issued forecasts and forecast-window exogenous inputs enter after the origin; future USGS observations are reserved for scoring. Main CRPS uses GloFAS-supported days 1–28, while NWS is compared over common leads 1–8.

INFERENCE

Scalable Bayesian quantile fitting

Seven independent Dynamic Quantile Linear Model (DQLM) lanes target seven levels: 0.05, 0.20, 0.35, 0.50, 0.65, 0.80, and 0.95. The fitted quantile forecasts are checked for crossing in the application and combined into the posterior predictive distribution scored by CRPS.

7 quantile lanes

≤12,995 USGS archive days

~2–4 h fit + synthesis

Parallel quantile lanes

Each quantile level is fit as its own state-space model; the lanes can run concurrently before forecast quantiles are synthesized.

Posterior approximation

Mean-field variational Bayes (MFVB) retains closed-form Gaussian filtering and smoothing within latent-state blocks while approximating posterior dependence across blocks.

Non-conjugate updates


Variational Bayes with Laplace–Delta (VB-LD) uses local modes, curvature, and Delta-method moments for non-conjugate extended asymmetric Laplace (exAL) scale and skewness terms.

State evolution

Component discounting controls trend, multi-period seasonality, source corrections, and exogenous-covariate adjustments.

MODEL

One DQLM for source correction, then synthesis



DQLM correction stage. For a fixed quantile p_0 , one Dynamic Quantile Linear Model spans the historical archive and forecast window: USGS observations and retrospectives identify the latent flow quantile and source-specific discrepancies before the forecast origin; issued ensembles and staged covariates propagate forecast-window quantiles. Seven quantile lanes are then synthesized and scored only against held-out USGS observations.

Expanded quantile state-space model

For each target quantile p_0 , let T denote the forecast origin/cutoff. The Dynamic Quantile Linear Model (DQLM) links observed flow, retrospective product histories, and issued forecast members through one latent river-flow quantile. Product-family discrepancies learn source corrections before the origin; the same source channels, with forecast-window covariates $\mathbf{x}_{T+k}$, then define corrected predictive quantiles after the origin.

Observed measurements: $y_t^o \sim \text{exAL}_{p_0}(\mathbf{F}_t' \boldsymbol{\theta}_t + \zeta_t, \sigma^o, \gamma^o), \quad 1 \leq t \leq T,$

Retrospective products: $z_t^j \sim \text{exAL}_{p_0}(\mathbf{F}_t'(\boldsymbol{\theta}_t + \boldsymbol{\delta}_t^j) + \zeta_t, \sigma^j, \gamma^j), \quad j \in \mathcal{J}, \quad 1 \leq t \leq T,$

Forecast products: $y_t^{j,i}(k) \sim \text{exAL}_{p_0}(\mathbf{F}_{T+k}'(\boldsymbol{\theta}_{T+k} + \boldsymbol{\delta}_{T+k}^j + \zeta_{T+k}, \sigma^j, \gamma^j), \quad j \in \mathcal{J}, \quad i \in \mathcal{I}_j, \quad k \in \mathcal{K}_j,$

Trend and seasonal states: $\boldsymbol{\theta}_t | \boldsymbol{\theta}_{t-1} \sim \mathcal{N}(\mathbf{G}_t \boldsymbol{\theta}_{t-1}, \mathbf{W}_t^{\boldsymbol{\theta}}), \quad 1 \leq t \leq T + K_{\max},$

Source correction states: $\boldsymbol{\delta}_t^j | \boldsymbol{\delta}_{t-1}^j \sim \mathcal{N}(\mathbf{G}_t \boldsymbol{\delta}_{t-1}^j, \mathbf{W}_t^{\boldsymbol{\delta}^j}), \quad j \in \mathcal{J}, \quad 1 \leq t \leq T + K_{\max},$


Exogenous adjustment: $\zeta_t | \zeta_{t-1}, \boldsymbol{\psi}_{t-1} \sim \mathcal{N}(\lambda \zeta_{t-1} + \mathbf{x}_t' \boldsymbol{\psi}_{t-1}, w_t^{\zeta}), \quad 1 \leq t \leq T + K_{\max},$

Covariate-adjustment states: $\boldsymbol{\psi}_t | \boldsymbol{\psi}_{t-1} \sim \mathcal{N}(\boldsymbol{\psi}_{t-1}, \mathbf{W}_t^{\boldsymbol{\psi}}), \quad 1 \leq t \leq T + K_{\max}.$

exAL_{p_0} is the extended asymmetric Laplace working likelihood for target quantile p_0 ; σ and γ are source-specific scale and skewness terms. T is the forecast origin, j a product family, i an ensemble member, k a lead, and $K_{\max}$ the longest horizon. The shared quantile state is $\boldsymbol{\theta}_t$; source discrepancies/corrections are $\boldsymbol{\delta}_t^j$; and $(\zeta_t, \boldsymbol{\psi}_t)$ is the exogenous-adjustment block driven by covariates $\mathbf{x}_t$ such as precipitation, soil moisture, and the GDPC climate-index summary. Component discounting controls the evolution variances.


MAIN EVIDENCE

28-day CRPS comparison




FIT

Wet-period quantile fit



EXAMPLE


Two forecast-window examples



Origin-level illustrations. Two selected forecast windows show how corrected quantile forecasts combine into a posterior predictive envelope after the cutoff. The panels are illustrative checks of behavior at individual origins; the aggregate comparison remains the five-origin CRPS benchmark.

DYNAMICS

80-month latent component



Retrospective diagnostic. From the selected model at the 2022-12-25 origin. Curves show the 5th, 50th, and 95th target-quantile components over 2000–2022; shaded intervals mark dry and wet regimes.

INTERPRETATION

Main takeaway

Answer to the poster question. Historical product archives can be used to learn source-specific corrections before forecast products are blended. In this five-origin San Lorenzo study, the selected source-aware model leads the 28-day CRPS comparison at four origins, while the fitted quantile diagnostics show interpretable changes in spread across wet high-flow periods. The design is a targeted rolling-origin stress test, not a continuous hindcast.

